

بِسْمِ اللَّهِ الرَّحْمَنِ الرَّحِيمِ

Imports

```
!pip install ipython-autotime
```

```
Requirement already satisfied: ipython-autotime in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: ipython in /usr/local/lib/python3.6/dist-packages (from ipython-autotime)
Requirement already satisfied: traitlets>=4.2 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: pexpect; sys_platform != "win32" in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: simplegeneric>0.8 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: decorator in /usr/local/lib/python3.6/dist-packages (from ipython)
Requirement already satisfied: pickleshare in /usr/local/lib/python3.6/dist-packages (from ipython)
Requirement already satisfied: prompt-toolkit<2.0.0,>=1.0.4 in /usr/local/lib/python3.6/dist-packages (from ipython)
Requirement already satisfied: setuptools>=18.5 in /usr/local/lib/python3.6/dist-packages (from ipython)
Requirement already satisfied: pygments in /usr/local/lib/python3.6/dist-packages (from ipython)
Requirement already satisfied: ipython-genutils in /usr/local/lib/python3.6/dist-packages (from ipython)
Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from traitlets)
Requirement already satisfied: ptyprocess>=0.5 in /usr/local/lib/python3.6/dist-packages (from pexpect)
Requirement already satisfied: wcwidth in /usr/local/lib/python3.6/dist-packages (from prompt-toolkit)
```

```
# necessary imports
import os
import cv2
import numpy as np
from imutils import paths
from sklearn.preprocessing import LabelBinarizer
from tqdm import tqdm

import matplotlib.pyplot as plt
%matplotlib inline

from google.colab.patches import cv2_imshow

%load_ext autotime

📄 time: 120 µs (started: 2021-01-07 06:35:34 +00:00)

img_width = 150
img_height = 150

time: 724 µs (started: 2021-01-07 06:35:34 +00:00)
```

▼ InceptionResNetV2 Model

```
from keras.applications import InceptionResNetV2
from keras.models import Model
from keras.layers import Dense
from keras.layers import Flatten
```

time: 1.32 s (started: 2021-01-07 06:35:35 +00:00)

```
model = InceptionResNetV2(include_top=False, weights='imagenet', input_shape=(img_width, img_
```

time: 5.13 s (started: 2021-01-07 06:35:36 +00:00)

```
# add new classification layers
flat1 = Flatten()(model.layers[-1].output) # flatten last layer
class1 = Dense(1024, activation='relu')(flat1) # add FC layer on previous layer
class2 = Dense(1024, activation='relu')(class1) # add FC layer on previous layer
output = Dense(6, activation='softmax')(class2) # add softmax layer
```

time: 27.1 ms (started: 2021-01-07 06:35:41 +00:00)

```
# define the new model
model = Model(inputs=model.inputs, outputs=output)
model.summary()
```

block8_9_conv (Conv2D)	(None, 3, 3, 2080)	0	block8_9_conv[0][0]
block8_9 (Lambda)	(None, 3, 3, 2080)	0	block8_9_ac[0][0]
block8_9_ac (Activation)	(None, 3, 3, 2080)	0	block8_9_conv[0][0]
conv2d_200 (Conv2D)	(None, 3, 3, 192)	399360	block8_9_ac[0][0]
batch_normalization_200 (BatchN	(None, 3, 3, 192)	576	conv2d_200[0][0]
activation_200 (Activation)	(None, 3, 3, 192)	0	batch_normalization_200[0][0]
conv2d_201 (Conv2D)	(None, 3, 3, 224)	129024	activation_200[0][0]
batch_normalization_201 (BatchN	(None, 3, 3, 224)	672	conv2d_201[0][0]
activation_201 (Activation)	(None, 3, 3, 224)	0	batch_normalization_201[0][0]
conv2d_199 (Conv2D)	(None, 3, 3, 192)	399360	block8_9_ac[0][0]
conv2d_202 (Conv2D)	(None, 3, 3, 256)	172032	activation_201[0][0]
batch_normalization_199 (BatchN	(None, 3, 3, 192)	576	conv2d_199[0][0]

batch_normalization_202 (BatchN	(None, 3, 3, 256)	768	conv2d_202[0][0]
activation_199 (Activation)	(None, 3, 3, 192)	0	batch_normalization_
activation_202 (Activation)	(None, 3, 3, 256)	0	batch_normalization_
block8_10_mixed (Concatenate)	(None, 3, 3, 448)	0	activation_199[0][0] activation_202[0][0]
block8_10_conv (Conv2D)	(None, 3, 3, 2080)	933920	block8_10_mixed[0][0]
block8_10 (Lambda)	(None, 3, 3, 2080)	0	block8_9_ac[0][0] block8_10_conv[0][0]
conv_7b (Conv2D)	(None, 3, 3, 1536)	3194880	block8_10[0][0]
conv_7b_bn (BatchNormalization)	(None, 3, 3, 1536)	4608	conv_7b[0][0]
conv_7b_ac (Activation)	(None, 3, 3, 1536)	0	conv_7b_bn[0][0]
flatten (Flatten)	(None, 13824)	0	conv_7b_ac[0][0]
dense (Dense)	(None, 1024)	14156800	flatten[0][0]
dense_1 (Dense)	(None, 1024)	1049600	dense[0][0]
dense_2 (Dense)	(None, 6)	6150	dense_1[0][0]

=====

Total params: 69,549,286
 Trainable params: 69,488,742
 Non-trainable params: 60,544

time: 275 ms (started: 2021-01-07 06:35:41 +00:00)

▼ Loading Data

```
# !unzip "/content/drive/MyDrive/CV/Assignment 3/intel-image-classification.zip" -d "/content
```

```
time: 680 µs (started: 2021-01-07 06:35:41 +00:00)
```

```
# !unzip "/content/drive/MyDrive/CV/Assignment 3/Test_data.zip" -d "/content/drive/MyDrive/CV
```

```
time: 436 µs (started: 2021-01-07 06:35:41 +00:00)
```

```
# A function to load data from a given directory
```

```
def load_data(data_dir):
    data = []
    labels = []
    class_dirs = os.listdir(data_dir)
```

```

for direc in class_dirs:
    # i=0
    class_dir = os.path.join(data_dir, direc)
    for imagepath in tqdm(list(paths.list_images(class_dir))):
        image = cv2.imread(imagepath)
        image = cv2.resize(image, (img_width, img_height)) # incase images not of same size
        data.append(image)
        labels.append(direc)
    # i = i+1
    # if (i==10):
    #     break
# normalizing and converting to numpy array format
data = np.array(data, dtype='float')/255.0
labels = np.array(labels)
return data, labels

```

time: 9.32 ms (started: 2021-01-07 06:35:41 +00:00)

```

train_dir = "/content/drive/MyDrive/CV/Assignment 3/seg_train/seg_train/"
test_dir = "/content/drive/MyDrive/CV/Assignment 3/seg_test/seg_test/"
pred_dir = "/content/drive/MyDrive/CV/Assignment 3/pred/seg_pred/seg_pred/"

```

time: 770 µs (started: 2021-01-07 06:35:41 +00:00)

Compile the model

```

import tensorflow as tf
import keras
from keras.optimizers import RMSprop

```

time: 718 µs (started: 2021-01-07 06:35:41 +00:00)

```

initial_learning_rate = 0.1
lr_schedule = tf.keras.optimizers.schedules.ExponentialDecay(
    initial_learning_rate,
    decay_steps=7000,
    decay_rate=0.96,
    staircase=True)

opt = keras.optimizers.RMSprop(learning_rate=lr_schedule)
model.compile(loss='categorical_crossentropy', optimizer=opt, metrics=['accuracy'])

```

time: 38 ms (started: 2021-01-07 06:35:41 +00:00)

```

# from keras.optimizers import SGD
# sgd = SGD(lr=0.001, decay=1e-7, momentum=.9)
# model.compile(loss='categorical_crossentropy',
#               optimizer=sgd,

```

```
# metrics=['accuracy'])

time: 634 µs (started: 2021-01-07 06:35:41 +00:00)

from keras.preprocessing.image import ImageDataGenerator

time: 7.17 ms (started: 2021-01-07 06:35:41 +00:00)

img_height = 150
img_width = 150
batch_size = 16
nb_epochs = 10

time: 715 µs (started: 2021-01-07 06:35:41 +00:00)

train_datagen = ImageDataGenerator(
    rescale=1./255,
    shear_range=0.2,
    zoom_range=0.2, # zoom
    rotation_range=10, # rotation
    width_shift_range=0.2, # horizontal shift
    height_shift_range=0.2, # vertical shift
    horizontal_flip=True) # horizontal flip
    # channel_shift_range = [-0.1, 0.1]
    # )
    # ,validation_split=0.3) # set validation split

time: 1.78 ms (started: 2021-01-07 06:35:41 +00:00)

train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(img_height, img_width),
    batch_size=batch_size,
    shuffle=True,
    class_mode='categorical',
    interpolation="nearest")
    # subset='training') # set as training data

Found 14034 images belonging to 6 classes.
time: 343 ms (started: 2021-01-07 06:35:41 +00:00)

val_datagen = ImageDataGenerator(rescale=1. / 255)

time: 2.43 ms (started: 2021-01-07 06:35:42 +00:00)

validation_generator = val_datagen.flow_from_directory(
    test_dir, # directory for validation data
    target_size=(img_height, img_width),
    batch_size=batch_size,
```

```
class_mode='categorical')
# subset='validation') # set as validation data


Found 3000 images belonging to 6 classes.
time: 113 ms (started: 2021-01-07 06:35:42 +00:00)
```

```
# validation_generator[0]
```

```
time: 1.3 ms (started: 2021-01-07 06:35:42 +00:00)
```

```
H = model.fit(
    train_generator,
    steps_per_epoch = train_generator.samples // batch_size,
    validation_data = validation_generator,
    validation_steps = validation_generator.samples // batch_size,
    epochs = nb_epochs)
```

```
Epoch 1/10
877/877 [=====] - 3905s 4s/step - loss: 22897.4580 - accuracy:
Epoch 2/10
877/877 [=====] - 164s 187ms/step - loss: 1.8664 - accuracy: 0
Epoch 3/10
877/877 [=====] - 166s 189ms/step - loss: 1.8042 - accuracy: 0
Epoch 4/10
877/877 [=====] - 166s 189ms/step - loss: 1.8070 - accuracy: 0
Epoch 5/10
877/877 [=====] - 167s 190ms/step - loss: 1.8064 - accuracy: 0
Epoch 6/10
877/877 [=====] - 166s 189ms/step - loss: 1.8063 - accuracy: 0
Epoch 7/10
877/877 [=====] - 166s 190ms/step - loss: 1.8057 - accuracy: 0
Epoch 8/10
877/877 [=====] - 166s 189ms/step - loss: 1.8044 - accuracy: 0
Epoch 9/10
877/877 [=====] - 165s 188ms/step - loss: 1.8051 - accuracy: 0
Epoch 10/10
877/877 [=====] - 166s 189ms/step - loss: 1.8051 - accuracy: 0
time: 1h 29min 58s (started: 2021-01-07 06:35:42 +00:00)
```



```
save_path = '/content/drive/MyDrive/CV/Assignment 3/InceptionResNetV2_Aug_Adam'
```

```
time: 1.19 ms (started: 2021-01-07 08:05:40 +00:00)
```

```
# save the model's trained weights
model.save_weights(save_path+"transfer_trained_wts.h5")
```


```
time: 1.95 s (started: 2021-01-07 08:14:26 +00:00)
```

```
# model.load_weights('/content/drive/MyDrive/CV/Assignment 3/vgg_aug_transfer_trained_wts.h5')
```

time: 680 μ s (started: 2021-01-07 08:14:29 +00:00)

```
simple_acc = H.history['accuracy']
plt.plot([1 - acc for acc in simple_acc])
plt.title('Error for a InceptionResNetV2 model using augmentation with Adam optimizer & adapt
plt.ylabel('Error')
plt.xlabel('Epoch')
plt.savefig(save_path+'/simple_acc_error.png')
plt.show()
```

Error for a InceptionResNetV2 model using augmentation with Adam optimizer & adaptive learning rate



time: 184 ms (started: 2021-01-07 08:15:20 +00:00)


```
simple_loss = H.history['loss']
plt.plot([los for los in simple_loss])
plt.title('Loss for a InceptionResNetV2_Aug model using augmentation with Adam optimizer & ad
plt.ylabel('Error')
plt.xlabel('Epoch')
plt.savefig(save_path+'/simple_loss_error.png')
plt.show()
```

```

simple_val_loss = H.history['val_loss']
plt.plot([los for los in simple_val_loss])
plt.title('Validation Loss for a InceptionResNetV2_Aug model using augmentation with Adam opt
plt.ylabel('Validation error')
plt.xlabel('Epoch')
plt.savefig(save_path+'/simple_Validation_loss_error.png')
plt.show()

```

Validation Loss for a InceptionResNetV2_Aug model using augmentation with Adam optimizer & adaptive learning rate




time: 198 ms (started: 2021-01-07 08:15:24 +00:00)

```

simple_val_acc = H.history['val_accuracy']
plt.plot([1 - acc for acc in simple_val_acc])
plt.title('Validation error for a InceptionResNetV2_Aug model using augmentation with Adam op
plt.ylabel('Validation error')
plt.xlabel('Epoch')
plt.savefig(save_path+'/simple_Validation_error.png')
plt.show()

```

Validation error for a InceptionResNetV2_Aug model using augmentation with Adam optimizer & adaptive learning rate



time: 196 ms (started: 2021-01-07 08:15:25 +00:00)


```
print('loading pred images')
X_test, y_test = load_data(pred_dir)

loading pred images
100%|██████████| 1128/1128 [08:34<00:00, 2.19it/s]
100%|██████████| 1297/1297 [10:04<00:00, 2.15it/s]
100%|██████████| 1330/1330 [10:14<00:00, 2.17it/s]
100%|██████████| 1166/1166 [09:06<00:00, 2.13it/s]
100%|██████████| 1144/1144 [09:30<00:00, 2.01it/s]
100%|██████████| 1236/1236 [09:43<00:00, 2.12it/s]
time: 57min 23s (started: 2021-01-07 08:15:25 +00:00)
```

```
lb = LabelBinarizer()
y_test = lb.fit_transform(y_test)

time: 11.8 ms (started: 2021-01-07 09:12:49 +00:00)
```

```
score = model.evaluate(X_test, y_test, batch_size=64)
print('Test Loss = ', score[0])
print('Test Accuracy = ', score[1])

115/115 [=====] - 20s 144ms/step - loss: 368.8753 - accuracy: 0.1606629192829132
Test Loss = 368.8753356933594
Test Accuracy = 0.1606629192829132
time: 22.1 s (started: 2021-01-07 09:12:49 +00:00)
```

```
'''CONFUSION MATRIX'''
# Making prediction
y_pred = model.predict(X_test)
y_true = np.argmax(y_test, axis=-1)

# Plotting the confusion matrix
from sklearn.metrics import confusion_matrix
confusion_mtx = confusion_matrix(y_true, np.argmax(y_pred, axis=1))

time: 21.8 s (started: 2021-01-07 09:13:11 +00:00)
```

```
mask = y_pred==y_test
correct = np.count_nonzero(mask)
print (correct*100.0/y_pred.size)
result = y_pred.astype(int)


1.1185682326621924
time: 5.16 ms (started: 2021-01-07 09:13:33 +00:00)
```

```
confusion_mtx

array([[ 0, 1125,  8,  0,  0, 11],
       [ 0, 1157,  3,  0,  0,  6],
```

```
[ 0, 1305, 12, 0, 0, 13],
[ 0, 1267, 15, 0, 0, 15],
[ 0, 1101, 14, 0, 0, 13],
[ 0, 1225, 7, 0, 0, 4]]time: 2.58 ms (started: 2021-01-07 09:13:33)
```

```
def plot_confusion_matrix(df_confusion, title='Confusion matrix', cmap=plt.cm.YlOrRd):
    plt.matshow(df_confusion, cmap=cmap) # imshow
    plt.colorbar()
    tick_marks = np.arange(6)
    names = ["Buildings", "Forest", "Glacier", "Mountain", "Sea", "Street"]
    plt.xticks(tick_marks, names, rotation=45)
    plt.yticks(tick_marks, names)
    plt.ylabel("Actual")
    plt.xlabel("Predicted")
#call function
plot_confusion_matrix(confusion_mtx)
```



time: 205 ms (started: 2021-01-07 09:13:33 +00:00)


```
import seaborn as sns

class_names = ['buildings', 'street', 'forest', 'glacier', 'mountain', 'sea']

class_names = sorted(class_names)

sns.heatmap(confusion_mtx, xticklabels=class_names, yticklabels=class_names,
            annot=True, fmt='d', cmap="YlGnBu")
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f6978e74320>



```
fig, axis = plt.subplots(1, 2, figsize=(20, 4))
```

```
axis[0].plot(H.history['accuracy'],
             label='Train accuracy with augmentation',
             c='tomato', ls='-')
axis[0].plot(H.history['val_accuracy'],
             label='Validation accuracy with augmentation',
             c='magenta', ls='-')
```

```
axis[0].set_xlabel('Epoch')
axis[0].set_ylabel('Accuracy')
axis[0].legend(loc='upper left')
```

```
axis[1].plot(H.history['loss'],
             label='Train loss with augmentation',
             c='tomato', ls='-')
axis[1].plot(H.history['val_loss'],
             label='Validation loss with augmentation',
             c='magenta', ls='-')
```

```
axis[1].set_xlabel('Epoch')
axis[1].set_ylabel('loss')
axis[1].legend(loc='upper left')
plt.savefig(save_path+'/simple_Validation_error&loss.png')
plt.show()
```

File "<ipython-input-42-5fcda59f6d64>", line 27

```
plt.show()
```

^

SyntaxError: invalid syntax

SEARCH STACK OVERFLOW

```
def visualize_data(images, categories, class_names):
    fig = plt.figure(figsize=(14, 6))
    fig.patch.set_facecolor('white')
    for i in range(3 * 6):
        plt.subplot(3, 6, i+1)
        plt.xticks([])
```

```
plt.yticks([])
plt.imshow(images[i])
class_index = categories[i].argmax()
plt.xlabel(class_names[class_index])
plt.show()
```

```
visualize_data(X_test*255, y_test, names)
```