

بِسْمِ اللَّهِ الرَّحْمَنِ الرَّحِيمِ

Imports

```
!pip install ipython-autotime
```

```
Requirement already satisfied: ipython-autotime in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: ipython in /usr/local/lib/python3.6/dist-packages (from ipython-autotime)
Requirement already satisfied: traitlets>=4.2 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: simplegeneric>0.8 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: setuptools>=18.5 in /usr/local/lib/python3.6/dist-packages
Requirement already satisfied: pickleshare in /usr/local/lib/python3.6/dist-packages (from ipython)
Requirement already satisfied: decorator in /usr/local/lib/python3.6/dist-packages (from ipython)
Requirement already satisfied: pygments in /usr/local/lib/python3.6/dist-packages (from ipython)
Requirement already satisfied: prompt-toolkit<2.0.0,>=1.0.4 in /usr/local/lib/python3.6/dist-packages (from ipython)
Requirement already satisfied: pexpect; sys_platform != "win32" in /usr/local/lib/python3.6/dist-packages (from ipython)
Requirement already satisfied: ipython-genutils in /usr/local/lib/python3.6/dist-packages (from ipython)
Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from traitlets)
Requirement already satisfied: wcwidth in /usr/local/lib/python3.6/dist-packages (from prompt-toolkit)
Requirement already satisfied: ptyprocess>=0.5 in /usr/local/lib/python3.6/dist-packages (from pexpect)
```

```
# necessary imports
import os
import cv2
import numpy as np
from imutils import paths
from sklearn.preprocessing import LabelBinarizer
from tqdm import tqdm

import matplotlib.pyplot as plt
%matplotlib inline

from google.colab.patches import cv2_imshow

%load_ext autotime

time: 109 µs (started: 2021-01-05 21:01:52 +00:00)

img_width = 150
img_height = 150

time: 1.19 ms (started: 2021-01-05 21:01:52 +00:00)
```

▼ VGG Model

```
from keras.applications.vgg16 import VGG16
from keras.models import Model
from keras.layers import Dense
from keras.layers import Flatten
```

```
time: 3.05 s (started: 2021-01-05 21:01:52 +00:00)
```

```
# load VGG16 model without classification layers
model = VGG16(include_top=False, input_shape=(150, 150, 3))
```

```
time: 2.67 s (started: 2021-01-05 21:01:55 +00:00)
```

```
# add new classification layers
flat1 = Flatten()(model.layers[-1].output) # flatten last layer
class1 = Dense(1024, activation='relu')(flat1) # add FC layer on previous layer
output = Dense(6, activation='softmax')(class1) # add softmax layer
```

```
time: 21 ms (started: 2021-01-05 21:01:58 +00:00)
```

```
# define the new model
model = Model(inputs=model.inputs, outputs=output)
model.summary()
```

```
Model: "model"
```

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 150, 150, 3)]	0
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168
block3_conv2 (Conv2D)	(None, 37, 37, 256)	590080
block3_conv3 (Conv2D)	(None, 37, 37, 256)	590080

block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
block4_conv1 (Conv2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Conv2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Conv2D)	(None, 18, 18, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
block5_conv1 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 9, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
flatten (Flatten)	(None, 8192)	0
dense (Dense)	(None, 1024)	8389632
dense_1 (Dense)	(None, 6)	6150

=====

Total params: 23,110,470
 Trainable params: 23,110,470
 Non-trainable params: 0

time: 11.7 ms (started: 2021-01-05 21:01:58 +00:00)

▼ Loading Data

```
# !unzip "/content/drive/MyDrive/CV/Assignment 3/intel-image-classification.zip" -d "/content"
time: 940 µs (started: 2021-01-05 21:01:58 +00:00)
```

```
# !unzip "/content/drive/MyDrive/CV/Assignment 3/Test_data.zip" -d "/content/drive/MyDrive/CV"
time: 574 µs (started: 2021-01-05 21:01:58 +00:00)
```

```
# A function to load data from a given directory
def load_data(data_dir):
    data = []
    labels = []
    class_dirs = os.listdir(data_dir)

    for direc in class_dirs:
        # i=0
        class_dir = os.path.join(data_dir, direc)
        for imagepath in tqdm(list(paths.list_images(class_dir))):
            image = cv2.imread(imagepath)
```

```

image = cv2.imread(imagepath)
image = cv2.resize(image, (img_width, img_height)) # incase images not of same size
data.append(image)
labels.append(direc)
# i = i+1
# if (i==10):
#     break
# normalizing and converting to numpy array format
data = np.array(data, dtype='float')/255.0
labels = np.array(labels)
return data, labels

```

time: 12 ms (started: 2021-01-05 21:01:58 +00:00)

```

train_dir = "/content/drive/MyDrive/CV/Assignment 3/seg_train/seg_train/"
test_dir = "/content/drive/MyDrive/CV/Assignment 3/seg_test/seg_test/"
pred_dir = "/content/drive/MyDrive/CV/Assignment 3/pred/seg_pred/seg_pred/"

```

time: 1.16 ms (started: 2021-01-05 21:01:58 +00:00)

Compile the model

```

from keras.optimizers import SGD
sgd = SGD(lr=0.001, decay=1e-7, momentum=.9)
model.compile(loss='categorical_crossentropy',
              optimizer=sgd,
              metrics=['accuracy'])

```

time: 21.8 ms (started: 2021-01-05 21:01:58 +00:00)

```

from keras.preprocessing.image import ImageDataGenerator

```

time: 2.52 ms (started: 2021-01-05 21:01:58 +00:00)

```

img_height = 150
img_width = 150
batch_size = 16
nb_epochs = 100

```

time: 3.57 ms (started: 2021-01-05 21:01:58 +00:00)

```

train_datagen = ImageDataGenerator(
    rescale=1./255,
    shear_range=0.2,
    zoom_range=0.2, # zoom
    rotation_range=10, # rotation
    width_shift_range=0.2, # horizontal shift
    height_shift_range=0.2, # vertical shift
    horizontal_flip=True) # horizontal flip

```

```

horizontal_flip=True, # horizontal flip
# channel_shift_range = [-0.1, 0.1]
# )
# ,validation_split=0.3) # set validation split

time: 2.83 ms (started: 2021-01-05 21:05:44 +00:00)

```

```

train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(img_height, img_width),
    batch_size=batch_size,
    shuffle=True,
    class_mode='categorical',
    interpolation="nearest")
# subset='training') # set as training data

Found 14033 images belonging to 6 classes.
time: 336 ms (started: 2021-01-05 21:05:47 +00:00)

```

```

val_datagen = ImageDataGenerator(rescale=1. / 255)

time: 805 µs (started: 2021-01-05 21:02:00 +00:00)

```

```

validation_generator = val_datagen.flow_from_directory(
    test_dir, # directory for validation data
    target_size=(img_height, img_width),
    batch_size=batch_size,
    class_mode='categorical')
# subset='validation') # set as validation data

Found 3000 images belonging to 6 classes.
time: 421 ms (started: 2021-01-05 21:02:00 +00:00)

```

```

# validation_generator[0]

time: 703 µs (started: 2021-01-05 21:02:00 +00:00)

```

```

H = model.fit(
    train_generator,
    steps_per_epoch = train_generator.samples // batch_size,
    validation_data = validation_generator,
    validation_steps = validation_generator.samples // batch_size,
    epochs = nb_epochs)

Epoch 73/100
877/877 [=====] - 130s 148ms/step - loss: 0.0127 - accuracy:
Epoch 74/100
877/877 [=====] - 130s 148ms/step - loss: 0.0113 - accuracy:
Epoch 75/100
877/877 [=====] - 129s 148ms/step - loss: 0.0119 - accuracy:
Epoch 76/100

```

```

Epoch 70/100
877/877 [=====] - 129s 147ms/step - loss: 0.0179 - accuracy:
Epoch 77/100
877/877 [=====] - 130s 148ms/step - loss: 0.0127 - accuracy:
Epoch 78/100
877/877 [=====] - 129s 148ms/step - loss: 0.0187 - accuracy:
Epoch 79/100
877/877 [=====] - 130s 148ms/step - loss: 0.0116 - accuracy:
Epoch 80/100
877/877 [=====] - 129s 147ms/step - loss: 0.0163 - accuracy:
Epoch 81/100
877/877 [=====] - 129s 147ms/step - loss: 0.0172 - accuracy:
Epoch 82/100
877/877 [=====] - 129s 148ms/step - loss: 0.0263 - accuracy:
Epoch 83/100
877/877 [=====] - 129s 147ms/step - loss: 0.0156 - accuracy:
Epoch 84/100
877/877 [=====] - 129s 147ms/step - loss: 0.0099 - accuracy:
Epoch 85/100
877/877 [=====] - 129s 147ms/step - loss: 0.0181 - accuracy:
Epoch 86/100
877/877 [=====] - 129s 148ms/step - loss: 0.0175 - accuracy:
Epoch 87/100
877/877 [=====] - 129s 147ms/step - loss: 0.0110 - accuracy:
Epoch 88/100
877/877 [=====] - 129s 147ms/step - loss: 0.0095 - accuracy:
Epoch 89/100

877/877 [=====] - 129s 148ms/step - loss: 0.0086 - accuracy:
Epoch 90/100
877/877 [=====] - 129s 148ms/step - loss: 0.0125 - accuracy:
Epoch 91/100
877/877 [=====] - 130s 148ms/step - loss: 0.0141 - accuracy:
Epoch 92/100
877/877 [=====] - 129s 147ms/step - loss: 0.0123 - accuracy:
Epoch 93/100
877/877 [=====] - 129s 148ms/step - loss: 0.0071 - accuracy:
Epoch 94/100
877/877 [=====] - 129s 148ms/step - loss: 0.0084 - accuracy:
Epoch 95/100
877/877 [=====] - 129s 148ms/step - loss: 0.0114 - accuracy:
Epoch 96/100
877/877 [=====] - 130s 148ms/step - loss: 0.0127 - accuracy:
Epoch 97/100
877/877 [=====] - 129s 148ms/step - loss: 0.0131 - accuracy:
Epoch 98/100
877/877 [=====] - 130s 148ms/step - loss: 0.0074 - accuracy:
Epoch 99/100
877/877 [=====] - 129s 147ms/step - loss: 0.0062 - accuracy:
Epoch 100/100
877/877 [=====] - 129s 147ms/step - loss: 0.0225 - accuracy:
time: 4h 8min 12s (started: 2021-01-05 21:06:29 +00:00)

```

```

# save the model's trained weights
model.save_weights('/content/drive/MyDrive/CV/Assignment 3/vgg_aug_transfer_trained_wts.h5')


```

time: 3.02 s (started: 2021-01-06 01:14:45 +00:00)

```
# model.load_weights('/content/drive/MyDrive/CV/Assignment 3/vgg_aug_transfer_trained_wts.h5')
```

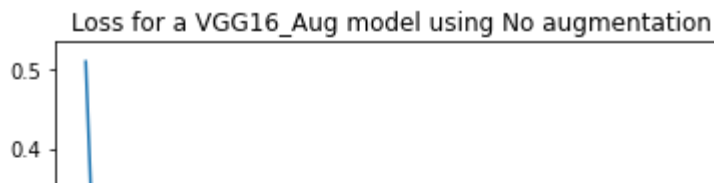
time: 1.35 ms (started: 2021-01-06 01:14:48 +00:00)

```
simple_acc = H.history['accuracy']
plt.plot([1 - acc for acc in simple_acc])
plt.title('Error for a VGG16_Aug model using No augmentation')
plt.ylabel('Error')
plt.xlabel('Epoch')
plt.savefig('/content/drive/MyDrive/CV/Assignment 3/VGG16_Aug/simple_acc_error.png')
plt.show()
```




time: 208 ms (started: 2021-01-06 01:14:48 +00:00)

```
simple_loss = H.history['loss']
plt.plot([los for los in simple_loss])
plt.title('Loss for a VGG16_Aug model using No augmentation')
plt.ylabel('Error')
plt.xlabel('Epoch')
plt.savefig('/content/drive/MyDrive/CV/Assignment 3/VGG16_Aug/simple_loss_error.png')
plt.show()
```




```
simple_val_loss = H.history['val_loss']
plt.plot([los for los in simple_val_loss])
plt.title('Validation Loss for a VGG16_Aug model using No augmentation')
plt.ylabel('Validation error')
plt.xlabel('Epoch')
plt.savefig('/content/drive/MyDrive/CV/Assignment 3/VGG16_Aug/simple_Validation_loss_error.png')
plt.show()
```



time: 197 ms (started: 2021-01-06 01:14:48 +00:00)

```
simple_val_acc = H.history['val_accuracy']
plt.plot([1 - acc for acc in simple_val_acc])
plt.title('Validation error for a VGG16_Aug model using No augmentation')
plt.ylabel('Validation error')
plt.xlabel('Epoch')
plt.savefig('/content/drive/MyDrive/CV/Assignment 3/VGG16_Aug/simple_Validation_error.png')
plt.show()
```


Validation error for a VGG16_Aug model using No augmentation



```
# test_datagen = ImageDataGenerator()
```

```
time: 510 µs (started: 2021-01-06 01:14:48 +00:00)
```

```
2021-01-06 01:14:48 +00:00
```

```
# test_generator = test_datagen.flow_from_directory(
#     pred_dir, # directory for prediction data
#     target_size=(img_height, img_width),
#     batch_size=batch_size,
#     class_mode='categorical',
#     subset='Prediction') # set as validation data
```

```
print('loading pred images')
```

```
X_test, y_test = load_data(pred_dir)
```

```
2%|| 33/1330 [00:00<00:04, 322.29it/s]loading pred images
100%|| 1330/1330 [00:04<00:00, 323.82it/s]
100%|| 1297/1297 [00:03<00:00, 361.66it/s]
100%|| 1128/1128 [00:03<00:00, 324.01it/s]
100%|| 1166/1166 [00:03<00:00, 336.23it/s]
100%|| 1236/1236 [00:03<00:00, 345.48it/s]
100%|| 1144/1144 [00:03<00:00, 339.00it/s]
time: 23.2 s (started: 2021-01-06 01:14:48 +00:00)
```

```
lb = LabelBinarizer()
```

```
y_test = lb.fit_transform(y_test)
```

```
time: 13.7 ms (started: 2021-01-06 01:15:12 +00:00)
```

```
score = model.evaluate(X_test, y_test, batch_size=64)
```

```
print('Test Loss = ', score[0])
```

```
print('Test Accuracy = ', score[1])
```

```
115/115 [=====] - 16s 126ms/step - loss: 1.0472 - accuracy: 0.8371
Test Loss = 1.0471714735031128
Test Accuracy = 0.8371455669403076
time: 17.2 s (started: 2021-01-06 01:15:12 +00:00)
```

```
'''CONFUSION MATRIX'''
```

```
# Making prediction
```

```
y_pred = model.predict(X_test)
```

```
y_true = np.argmax(y_test, axis=-1)
```

```
# Plotting the confusion matrix
```

```
from sklearn.metrics import confusion_matrix
```

```
confusion_mtx = confusion_matrix(y_true, np.argmax(y_pred, axis=1))
```

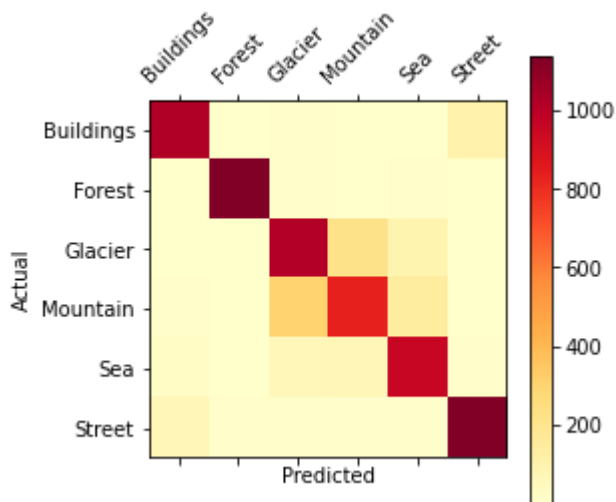
time: 18.5 s (started: 2021-01-06 01:15:29 +00:00)

```
confusion_mtx
```

```
array([[1022,    4,    8,    2,    3, 105],
       [   5, 1140,    8,    3,    8,    2],
       [   3,    2, 1017, 216,   91,    1],
       [   6,    4, 305, 834, 147,    1],
       [  21,    5,   59,   69, 962,   12],
       [  68,   11,    8,    6,    6, 1137]])
```



```
def plot_confusion_matrix(df_confusion, title='Confusion matrix', cmap=plt.cm.YlOrRd):
    plt.matshow(df_confusion, cmap=cmap) # imshow
    plt.colorbar()
    tick_marks = np.arange(6)
    names = ["Buildings", "Forest", "Glacier", "Mountain", "Sea", "Street"]
    plt.xticks(tick_marks, names, rotation=45)
    plt.yticks(tick_marks, names)
    plt.ylabel("Actual")
    plt.xlabel("Predicted")
#call function
plot_confusion_matrix(confusion_mtx)
```



time: 166 ms (started: 2021-01-06 01:15:47 +00:00)

```
import seaborn as sns
```


```
class_names = ['buildings','street','forest','glacier','mountain','sea']
```

```
class_names = sorted(class_names)
```

```
sns.heatmap(confusion_mtx, xticklabels=class_names, yticklabels=class_names,
```

```
annot=True, fmt='d', cmap="YlGnBu")
```


<matplotlib.axes._subplots.AxesSubplot at 0x7f37a223c5f8>



```
fig = plt.figure()
# fig.patch.set_facecolor('white')

plt.plot(H.history['accuracy'],
         label='train accuracy augmented',
         c='orange', ls='-')
plt.plot(H.history['val_accuracy'],
         label='validation accuracy augmented',
         c='blue', ls='-')


plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
plt.savefig('/content/drive/MyDrive/CV/Assignment 3/VGG16_Aug/Aug_Validation_error.png')
plt.show()
```



```
fig = plt.figure()
```

```
plt.plot(H.history['loss'],
         label='train loss augmented',
         c='orange', ls='-')
plt.plot(H.history['val_loss'],
         label='validation loss augmented',
         c='blue', ls='-')

plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='upper left')
plt.savefig('/content/drive/MyDrive/CV/Assignment 3/VGG16_Aug/Aug_Validation_&loss.png')
plt.show()
```



time: 223 ms (started: 2021-01-06 01:29:23 +00:00)



```
fig, axis = plt.subplots(1, 2, figsize=(20, 4))

axis[0].plot(H.history['accuracy'],
             label='Train accuracy with augmentation',
             c='tomato', ls='-')
axis[0].plot(H.history['val_accuracy'],
             label='Validation accuracy with augmentation',
             c='magenta', ls='-')

axis[0].set_xlabel('Epoch')
axis[0].set_ylabel('Accuracy')
axis[0].legend(loc='lower right')

axis[1].plot(H.history['loss'],
             label='Train loss with augmentation',
             c='tomato', ls='-')
axis[1].plot(H.history['val_loss'],
             label='Validation loss with augmentation',
             c='magenta', ls='-')
```

```
axis[1].set_xlabel('Epoch')  
axis[1].set_ylabel('loss')  
axis[1].legend(loc='upper left')  
plt.savefig('/content/drive/MyDrive/CV/Assignment 3/VGG16_Aug/Aug_Validation_error&loss.png')  
plt.show()
```



time: 452 ms (started: 2021-01-06 01:28:34 +00:00)