

بِسْمِ اللَّهِ الرَّحْمَنِ الرَّحِيمِ

Imports

```
!pip install ipython-autotime
```

```
Collecting ipython-autotime
```

```
  Downloading https://files.pythonhosted.org/packages/d6/c5/013f5aa3b56c6d2c58634bc97977
```

```
Requirement already satisfied: ipython in /usr/local/lib/python3.6/dist-packages (from ipython-autotime)
```

```
Requirement already satisfied: pygments in /usr/local/lib/python3.6/dist-packages (from ipython)
```

```
Requirement already satisfied: simplegeneric>0.8 in /usr/local/lib/python3.6/dist-packages (from ipython)
```

```
Requirement already satisfied: prompt-toolkit<2.0.0,>=1.0.4 in /usr/local/lib/python3.6/dist-packages (from ipython)
```

```
Requirement already satisfied: setuptools>=18.5 in /usr/local/lib/python3.6/dist-packages (from ipython)
```

```
Requirement already satisfied: pickleshare in /usr/local/lib/python3.6/dist-packages (from ipython)
```

```
Requirement already satisfied: pexpect; sys_platform != "win32" in /usr/local/lib/python3.6/dist-packages (from ipython)
```

```
Requirement already satisfied: traitlets>=4.2 in /usr/local/lib/python3.6/dist-packages (from ipython)
```

```
Requirement already satisfied: decorator in /usr/local/lib/python3.6/dist-packages (from ipython)
```

```
Requirement already satisfied: wcwidth in /usr/local/lib/python3.6/dist-packages (from ipython)
```

```
Requirement already satisfied: six>=1.9.0 in /usr/local/lib/python3.6/dist-packages (from ipython)
```

```
Requirement already satisfied: ptyprocess>=0.5 in /usr/local/lib/python3.6/dist-packages (from ipython)
```

```
Requirement already satisfied: ipython-genutils in /usr/local/lib/python3.6/dist-packages (from ipython)
```

```
Installing collected packages: ipython-autotime
```

```
Successfully installed ipython-autotime-0.3.0
```

```
# necessary imports
import os
import cv2
import numpy as np
from imutils import paths
from sklearn.preprocessing import LabelBinarizer
from tqdm import tqdm
```

```
import matplotlib.pyplot as plt
%matplotlib inline
```

```
from google.colab.patches import cv2_imshow
```

```
%load_ext autotime
```

```
time: 102 µs (started: 2021-01-06 02:15:54 +00:00)
```

```
img_width = 150
img_height = 150
```

```
time: 1.94 ms (started: 2021-01-06 02:15:54 +00:00)
```


▼ InceptionResNetV2 Model

```
from keras.applications import InceptionResNetV2
from keras.models import Model
from keras.layers import Dense
from keras.layers import Flatten
```

time: 1.49 s (started: 2021-01-06 02:15:54 +00:00)

```
model = InceptionResNetV2(include_top=False, weights='imagenet', input_shape=(img_width, img_
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/inception_resnet_v2/219062272/219055592 [=====] - 1s 0us/step
time: 11.8 s (started: 2021-01-06 02:15:56 +00:00)



```
# add new classification layers
flat1 = Flatten()(model.layers[-1].output) # flatten last layer
class1 = Dense(1024, activation='relu')(flat1) # add FC layer on previous layer
class2 = Dense(1024, activation='relu')(class1) # add FC layer on previous layer
output = Dense(6, activation='softmax')(class2) # add softmax layer
```

time: 28.7 ms (started: 2021-01-06 02:16:08 +00:00)

```
# define the new model
model = Model(inputs=model.inputs, outputs=output)
model.summary()
```

block8_9_conv (Conv2D)	(None, 3, 3, 2080)	0	block8_8_ac[0][0] block8_9_conv[0][0]
block8_9_ac (Activation)	(None, 3, 3, 2080)	0	block8_9[0][0]
conv2d_200 (Conv2D)	(None, 3, 3, 192)	399360	block8_9_ac[0][0]
batch_normalization_200 (Batch Normalization)	(None, 3, 3, 192)	576	conv2d_200[0][0]
activation_200 (Activation)	(None, 3, 3, 192)	0	batch_normalization_200[0][0]
conv2d_201 (Conv2D)	(None, 3, 3, 224)	129024	activation_200[0][0]
batch_normalization_201 (Batch Normalization)	(None, 3, 3, 224)	672	conv2d_201[0][0]
activation_201 (Activation)	(None, 3, 3, 224)	0	batch_normalization_201[0][0]
conv2d_199 (Conv2D)	(None, 3, 3, 192)	399360	block8_9_ac[0][0]
conv2d_202 (Conv2D)	(None, 3, 3, 256)	173824	activation_201[0][0]

conv2d_202 (Conv2D)	(None, 3, 3, 256)	172032	activation_201[0][0]
batch_normalization_199 (BatchN	(None, 3, 3, 192)	576	conv2d_199[0][0]
batch_normalization_202 (BatchN	(None, 3, 3, 256)	768	conv2d_202[0][0]
activation_199 (Activation)	(None, 3, 3, 192)	0	batch_normalization_
activation_202 (Activation)	(None, 3, 3, 256)	0	batch_normalization_
block8_10_mixed (Concatenate)	(None, 3, 3, 448)	0	activation_199[0][0] activation_202[0][0]
block8_10_conv (Conv2D)	(None, 3, 3, 2080)	933920	block8_10_mixed[0][0]
block8_10 (Lambda)	(None, 3, 3, 2080)	0	block8_9_ac[0][0] block8_10_conv[0][0]
conv_7b (Conv2D)	(None, 3, 3, 1536)	3194880	block8_10[0][0]
conv_7b_bn (BatchNormalization)	(None, 3, 3, 1536)	4608	conv_7b[0][0]
conv_7b_ac (Activation)	(None, 3, 3, 1536)	0	conv_7b_bn[0][0]
flatten (Flatten)	(None, 13824)	0	conv_7b_ac[0][0]
dense (Dense)	(None, 1024)	14156800	flatten[0][0]
dense_1 (Dense)	(None, 1024)	1049600	dense[0][0]
dense_2 (Dense)	(None, 6)	6150	dense_1[0][0]

=====

Total params: 69,549,286
 Trainable params: 69,488,742
 Non-trainable params: 60,544

time: 295 ms (started: 2021-01-06 02:16:08 +00:00)

▼ Loading Data

```
# !unzip "/content/drive/MyDrive/CV/Assignment 3/intel-image-classification.zip" -d "/content
```

```
time: 554 µs (started: 2021-01-06 02:16:08 +00:00)
```

```
# !unzip "/content/drive/MyDrive/CV/Assignment 3/Test_data.zip" -d "/content/drive/MyDrive/CV
```

```
time: 555 µs (started: 2021-01-06 02:16:08 +00:00)
```

```
# A function to load data from a given directory
```

```
def load_data(data_dir):
```

```

data = []
labels = []
class_dirs = os.listdir(data_dir)

for direc in class_dirs:
    # i=0
    class_dir = os.path.join(data_dir, direc)
    for imagepath in tqdm(list(paths.list_images(class_dir))):
        image = cv2.imread(imagepath)
        image = cv2.resize(image, (img_width, img_height)) # incase images not of same size
        data.append(image)
        labels.append(direc)
    # i = i+1
    # if (i==10):
    #     break
# normalizing and converting to numpy array format
data = np.array(data, dtype='float')/255.0
labels = np.array(labels)
return data, labels

```

time: 19.6 ms (started: 2021-01-06 02:16:08 +00:00)

```

train_dir = "/content/drive/MyDrive/CV/Assignment 3/seg_train/seg_train/"
test_dir = "/content/drive/MyDrive/CV/Assignment 3/seg_test/seg_test/"
pred_dir = "/content/drive/MyDrive/CV/Assignment 3/pred/seg_pred/seg_pred/"

```

time: 792 µs (started: 2021-01-06 02:16:08 +00:00)

Compile the model

```

# initial_learning_rate = 0.1
# lr_schedule = tf.keras.optimizers.schedules.ExponentialDecay(
#     initial_learning_rate,
#     decay_steps=7000,
#     decay_rate=0.96,
#     staircase=True)

# opt = keras.optimizers.Adam(learning_rate=lr_schedule)
# model.compile(loss='categorical_crossentropy', optimizer=opt, metrics=['accuracy'])

```

time: 2.69 ms (started: 2021-01-06 02:16:08 +00:00)

```

from keras.optimizers import SGD
sgd = SGD(lr=0.001, decay=1e-7, momentum=.9)
model.compile(loss='categorical_crossentropy',
              optimizer=sgd,
              metrics=['accuracy'])

```

time: 36.1 ms (started: 2021-01-06 02:16:08 +00:00)

```
from keras.preprocessing.image import ImageDataGenerator
```

```
time: 2.75 ms (started: 2021-01-06 02:16:08 +00:00)
```

```
img_height = 150
```

```
img_width = 150
```

```
batch_size = 16
```

```
nb_epochs = 25
```

```
time: 1.59 ms (started: 2021-01-06 02:16:08 +00:00)
```

```
train_datagen = ImageDataGenerator(  
    rescale=1./255,  
    shear_range=0.2,  
    zoom_range=0.2, # zoom  
    rotation_range=10, # rotation  
    width_shift_range=0.2, # horizontal shift  
    height_shift_range=0.2, # vertical shift  
    horizontal_flip=True) # horizontal flip  
# channel_shift_range = [-0.1, 0.1]  
# )  
# ,validation_split=0.3) # set validation split
```

```
time: 3.53 ms (started: 2021-01-06 02:16:08 +00:00)
```

```
train_generator = train_datagen.flow_from_directory(  
    train_dir,  
    target_size=(img_height, img_width),  
    batch_size=batch_size,  
    shuffle=True,  
    class_mode='categorical',  
    interpolation="nearest")  
# subset='training') # set as training data
```

```
Found 14033 images belonging to 6 classes.
```

```
time: 24.9 s (started: 2021-01-06 02:20:08 +00:00)
```

```
val_datagen = ImageDataGenerator(rescale=1. / 255)
```

```
time: 2.08 ms (started: 2021-01-06 02:20:33 +00:00)
```

```
validation_generator = val_datagen.flow_from_directory(  
    test_dir, # directory for validation data  
    target_size=(img_height, img_width),  
    batch_size=batch_size,  
    class_mode='categorical')  
# subset='validation') # set as validation data
```

Found 3000 images belonging to 6 classes.
time: 3.99 s (started: 2021-01-06 02:20:33 +00:00)

```
# validation_generator[0]
```

time: 816 µs (started: 2021-01-06 02:20:37 +00:00)

```
H = model.fit(
    train_generator,
    steps_per_epoch = train_generator.samples // batch_size,
    validation_data = validation_generator,
    validation_steps = validation_generator.samples // batch_size,
    epochs = nb_epochs)

Epoch 1/25
877/877 [=====] - 6053s 7s/step - loss: 0.8461 - accuracy: 0.66
Epoch 2/25
877/877 [=====] - 149s 170ms/step - loss: 0.4194 - accuracy: 0
Epoch 3/25
877/877 [=====] - 151s 172ms/step - loss: 0.3348 - accuracy: 0
Epoch 4/25
877/877 [=====] - 152s 173ms/step - loss: 0.2888 - accuracy: 0
Epoch 5/25
877/877 [=====] - 152s 173ms/step - loss: 0.2700 - accuracy: 0
Epoch 6/25
877/877 [=====] - 152s 174ms/step - loss: 0.2338 - accuracy: 0
Epoch 7/25
877/877 [=====] - 153s 174ms/step - loss: 0.2293 - accuracy: 0
Epoch 8/25
877/877 [=====] - 152s 174ms/step - loss: 0.2032 - accuracy: 0
Epoch 9/25
877/877 [=====] - 153s 174ms/step - loss: 0.1838 - accuracy: 0
Epoch 10/25
877/877 [=====] - 151s 173ms/step - loss: 0.1754 - accuracy: 0
Epoch 11/25
877/877 [=====] - 149s 170ms/step - loss: 0.1643 - accuracy: 0
Epoch 12/25
877/877 [=====] - 149s 169ms/step - loss: 0.1492 - accuracy: 0
Epoch 13/25
877/877 [=====] - 150s 171ms/step - loss: 0.1461 - accuracy: 0
Epoch 14/25
877/877 [=====] - 149s 170ms/step - loss: 0.1412 - accuracy: 0
Epoch 15/25
877/877 [=====] - 149s 170ms/step - loss: 0.1365 - accuracy: 0
Epoch 16/25
877/877 [=====] - 149s 169ms/step - loss: 0.1215 - accuracy: 0
Epoch 17/25
877/877 [=====] - 148s 169ms/step - loss: 0.1270 - accuracy: 0
Epoch 18/25
877/877 [=====] - 148s 168ms/step - loss: 0.1186 - accuracy: 0
Epoch 19/25
877/877 [=====] - 147s 168ms/step - loss: 0.1157 - accuracy: 0
Epoch 20/25
```

```

877/877 [=====] - 147s 168ms/step - loss: 0.0962 - accuracy: 0
Epoch 21/25
877/877 [=====] - 147s 168ms/step - loss: 0.1024 - accuracy: 0
Epoch 22/25
877/877 [=====] - 146s 167ms/step - loss: 0.0811 - accuracy: 0
Epoch 23/25
877/877 [=====] - 145s 166ms/step - loss: 0.0767 - accuracy: 0
Epoch 24/25
877/877 [=====] - 146s 166ms/step - loss: 0.0844 - accuracy: 0
Epoch 25/25
877/877 [=====] - 146s 166ms/step - loss: 0.0862 - accuracy: 0
time: 2h 40min 39s (started: 2021-01-06 02:20:37 +00:00)

```



```

# save the model's trained weights
model.save_weights('/content/drive/MyDrive/CV/Assignment 3/InceptionResNetV2_Aug_transfer_tra

time: 1.84 s (started: 2021-01-06 05:01:17 +00:00)

```

```

# model.load_weights('/content/drive/MyDrive/CV/Assignment 3/vgg_aug_transfer_trained_wts.h5'

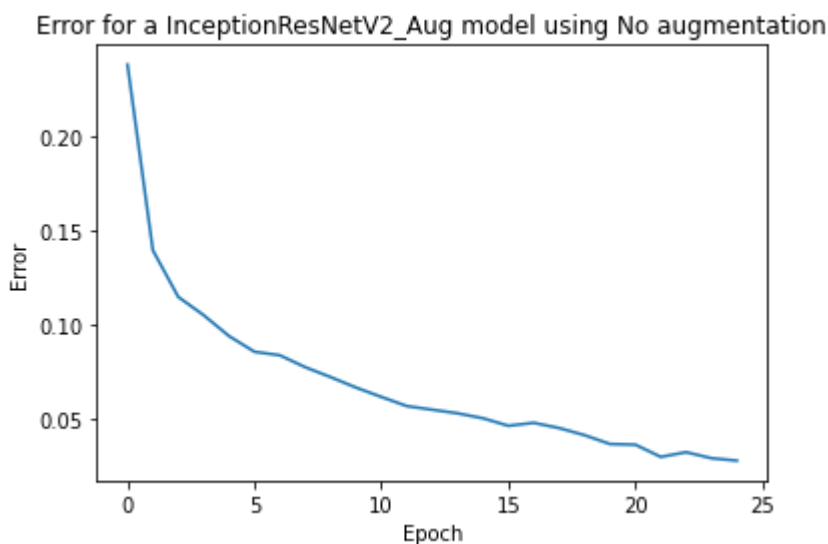
time: 610 µs (started: 2021-01-06 05:01:19 +00:00)

```

```

simple_acc = H.history['accuracy']
plt.plot([1 - acc for acc in simple_acc])
plt.title('Error for a InceptionResNetV2_Aug model using No augmentation')
plt.ylabel('Error')
plt.xlabel('Epoch')
plt.savefig('/content/drive/MyDrive/CV/Assignment 3/InceptionResNetV2_Aug/simple_acc_error.png')
plt.show()

```




```
time: 212 ms (started: 2021-01-06 05:01:19 +00:00)
```

```

simple_loss = H.history['loss']
plt.plot([los for los in simple_loss])
plt.title('Loss for a InceptionResNetV2 Aug model using No augmentation')


```

```
plt.ylabel('Error')
plt.xlabel('Epoch')
plt.savefig('/content/drive/MyDrive/CV/Assignment 3/InceptionResNetV2_Aug/simple_loss_error.p
plt.show()
```



time: 207 ms (started: 2021-01-06 05:01:19 +00:00)

```
simple_val_loss = H.history['val_loss']
plt.plot([los for los in simple_val_loss])
plt.title('Validation Loss for a InceptionResNetV2_Aug model using No augmentation')
plt.ylabel('Validation error')
plt.xlabel('Epoch')
plt.savefig('/content/drive/MyDrive/CV/Assignment 3/InceptionResNetV2_Aug/simple_Validation_1
plt.show()
```




time: 199 ms (started: 2021-01-06 05:01:19 +00:00)

```
simple_val_acc = H.history['val_accuracy']
plt.plot([1 - acc for acc in simple_val_acc])
plt.title('Validation error for a InceptionResNetV2_Aug model using No augmentation')
plt.ylabel('Validation error')
```



```
plt.xlabel('Epoch')
plt.savefig('/content/drive/MyDrive/CV/Assignment 3/InceptionResNetV2_Aug/simple_Validation_e
plt.show())
```

Validation error for a InceptionResNetV2_Aug model using No augmentation



time: 186 ms (started: 2021-01-06 05:01:19 +00:00)

```
# test_datagen = ImageDataGenerator()
```

time: 734 µs (started: 2021-01-06 05:01:19 +00:00)

```
# test_generator = test_datagen.flow_from_directory(
#     pred_dir, # directory for prediction data
#     target_size=(img_height, img_width),
#     batch_size=batch_size,
#     class_mode='categorical',
#     subset='Prediction') # set as validation data
```

time: 1.22 ms (started: 2021-01-06 05:01:20 +00:00)

```
print('loading pred images')
X_test, y_test = load_data(pred_dir)
```

```
loading pred images
100%|██████████| 1330/1330 [08:01<00:00, 2.76it/s]
100%|██████████| 1297/1297 [07:52<00:00, 2.75it/s]
100%|██████████| 1128/1128 [06:56<00:00, 2.71it/s]
100%|██████████| 1166/1166 [07:08<00:00, 2.72it/s]
100%|██████████| 1236/1236 [07:36<00:00, 2.70it/s]
100%|██████████| 1144/1144 [06:54<00:00, 2.76it/s]
time: 44min 48s (started: 2021-01-06 05:01:20 +00:00)
```

```
lb = LabelBinarizer()
y_test = lb.fit_transform(y_test)
```

time: 11.3 ms (started: 2021-01-06 05:46:08 +00:00)

```
score = model.evaluate(X_test, y_test, batch_size=64)
print('Test Loss = ', score[0])
print('Test Accuracy = ', score[1])
```

```
115/115 [=====] - 20s 143ms/step - loss: 0.9228 - accuracy: 0.82
Test Loss = 0.9227569103240967
Test Accuracy = 0.821668267250061
time: 22.2 s (started: 2021-01-06 05:46:08 +00:00)
```

```
'''CONFUSION MATRIX'''
# Making prediction
y_pred = model.predict(X_test)
y_true = np.argmax(y_test, axis=-1)


# Plotting the confusion matrix
from sklearn.metrics import confusion_matrix
confusion_mtx = confusion_matrix(y_true, np.argmax(y_pred, axis=1))

time: 21.9 s (started: 2021-01-06 05:46:30 +00:00)
```

confusion_mtx

```
array([[1051, 3, 11, 0, 8, 71],
       [ 5, 1137, 14, 1, 7, 2],
       [ 3, 8, 1161, 41, 112, 5],
       [ 3, 25, 489, 590, 188, 2],
       [ 15, 5, 106, 23, 968, 11],
       [ 106, 13, 17, 0, 8, 1092]])
time: 2.97 ms (started: 2021-01-06 05:46:51)
```

```
def plot_confusion_matrix(df_confusion, title='Confusion matrix', cmap=plt.cm.YlOrRd):
    plt.matshow(df_confusion, cmap=cmap) # imshow
    plt.colorbar()
    tick_marks = np.arange(6)
    names = ["Buildings", "Forest", "Glacier", "Mountain", "Sea", "Street"]
    plt.xticks(tick_marks, names, rotation=45)
    plt.yticks(tick_marks, names)
    plt.ylabel("Actual")
    plt.xlabel("Predicted")
#call function
plot_confusion_matrix(confusion_mtx)
```




```
import seaborn as sns
```

```
class_names = ['buildings','street','forest','glacier','mountain','sea']
```

```
class_names = sorted(class_names)
```

```
sns.heatmap(confusion_mtx, xticklabels=class_names, yticklabels=class_names,
            annot=True, fmt='d', cmap="YlGnBu")
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f78e4111e48>
```



```
time: 404 ms (started: 2021-01-06 05:46:52 +00:00)
```

```
fig, axis = plt.subplots(1, 2, figsize=(20, 4))
```

```
axis[0].plot(H.history['accuracy'],
            label='Train accuracy with augmentation',
            c='tomato', ls='-')
axis[0].plot(H.history['val_accuracy'],
            label='Validation accuracy with augmentation',
            c='magenta', ls='-')
```



```
axis[0].set_xlabel('Epoch')
axis[0].set_ylabel('Accuracy')
axis[0].legend(loc='upper left')
```

```

axis[1].plot(H.history['loss'],
             label='Train loss with augmentation',
             c='tomato', ls='-')
axis[1].plot(H.history['val_loss'],
             label='Validation loss with augmentation',
             c='magenta', ls='-')

axis[1].set_xlabel('Epoch')
axis[1].set_ylabel('loss')
axis[1].legend(loc='upper left')
plt.savefig('/content/drive/MyDrive/CV/Assignment 3/InceptionResNetV2_Aug/simple_Validation_e
plt.show()

```



time: 478 ms (started: 2021-01-06 05:46:53 +00:00)

